

论文分享: Tardis: Coverage-Guided Embedded Operating System Fuzzing

目录

1. 背景和挑战
2. 设计目标
3. 具体设计
4. 实验
5. 相关工作

背景和挑战

当前的基于覆盖率的fuzz方法，为什么不能直接用于embedded os？

- **Architecture:** Embedded OSs run on a variety of architectures and platforms
- **Intercafes:** fuzzer对OS的功能有一定依赖，很多Embedded OSs do not support the complete suite of interfaces required. Meanwhile, different Embedded OSs normally provide different interfaces for user land, thus a unified, OS-agnostic design is needed for coverage collection
- **Efficiency:** be efficient so that they would not significantly impact the execution throughput of the fuzzer

设计目标

- the coverage feedback mechanisms should be **software-based** and avoid relying on hardware-specific mechanisms to guarantee portability.
- these mechanisms should be **OS-agnostic**, i.e., it is able to operate without relying on any OS-specific features and can adapt to different OSes without additional costs.
- the mechanisms' implementations should be **highly efficient** and avoid introducing any significant overheads. 因为embedded OS对效率要求比较高

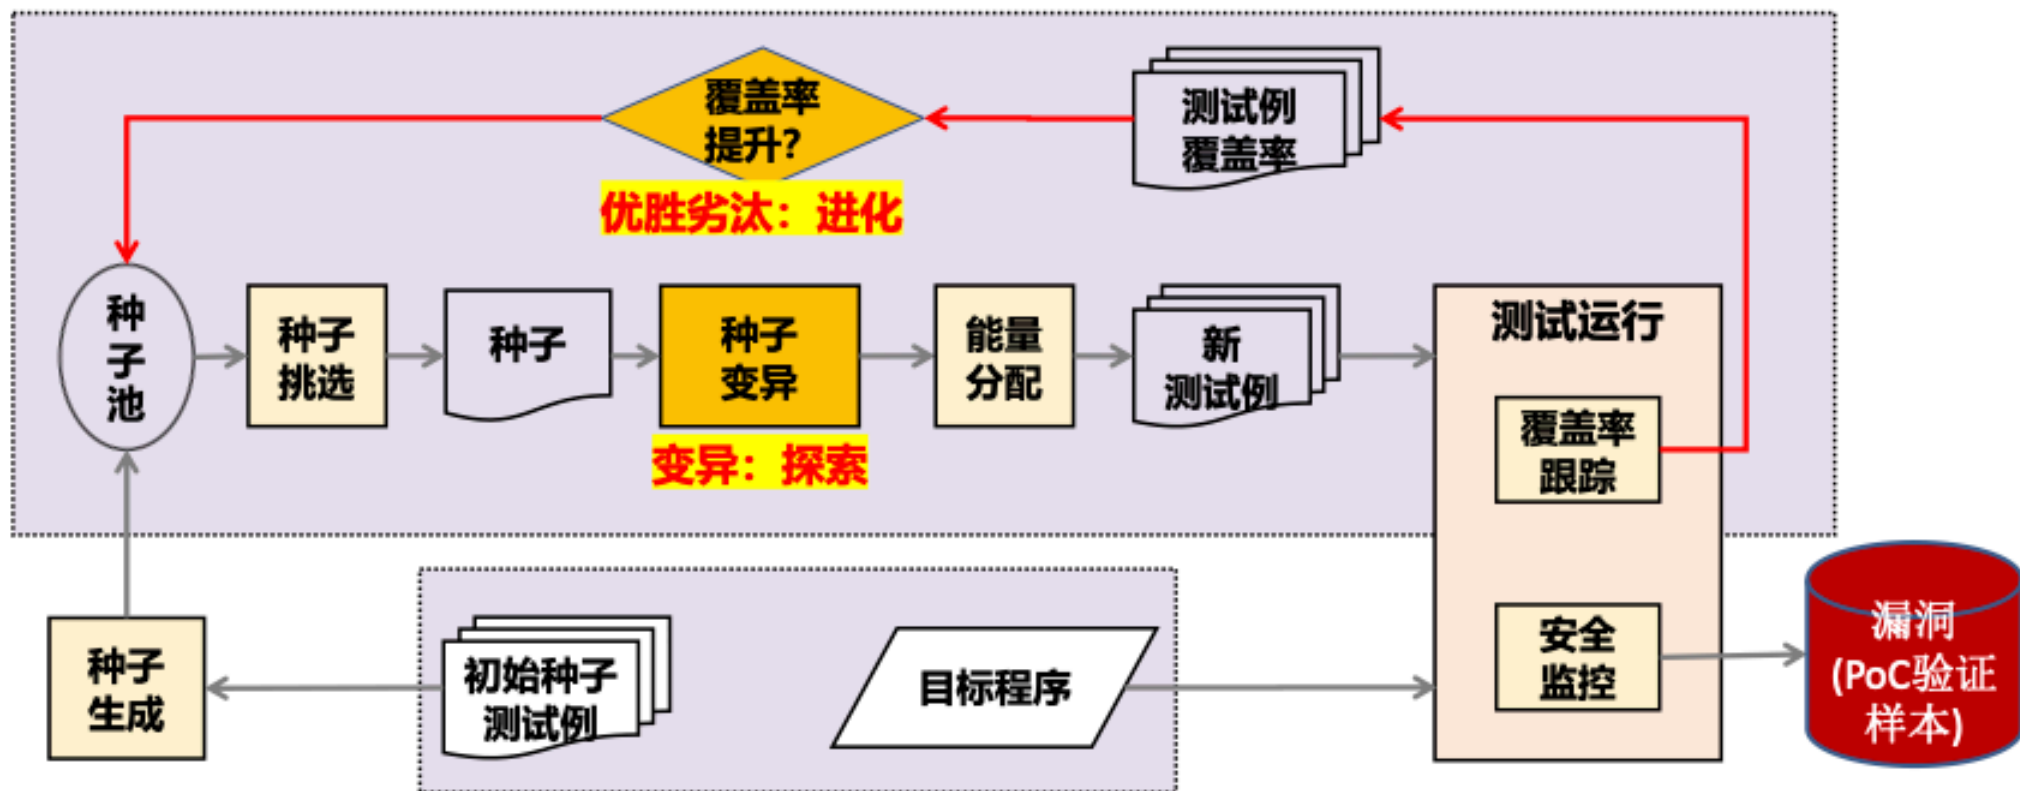
具体设计

第四代方案：智能模糊测试



清华大学
Tsinghua University

网络研究院
INSC



Google: 80%以上漏洞是通过fuzzing发现的 5

具体设计

两个关键的技术问题：

- 关于输入数据的自动调整算法，如何根据覆盖率反馈的结果，效率更高地生成触发新的执行路径的输入
- 覆盖率的部分，如何得到程序覆盖率，并且反馈给fuzzer，fuzzer再如何将分析结果交给上一步说的输入数据自动调整算法部分，用来指导后续的fuzz

具体设计

更加具体一点，想要收集程序覆盖率并且得到反馈，分为以下几个步骤：

1. coverage initialization。将被测试的代码划分为很多个基本块，并且通过回调函数，使得这个基本块如果被执行，那么就会调用另外一个用来统计基本块执行次数的函数
2. 覆盖率统计。不仅要知道这个基本块有没有被执行过，还希望知道这个执行路径有没有被执行过
3. coverage analyse。fuzzer希望能够以一个比较高的效率知道，相比之前的输入，是否有覆盖到新的执行路径

Coverage Initialization

现有的用户态fuzzing编译器插桩技术是如何做的：

- the Clang compiler supports a general instrumentation mechanism that can **inject user-defined callbacks** at the beginning of each code block. The callback **pc_trace_guard_init()** is used to initialize the whole collection procedure and **is called before entering program logic**, while the callback **pc_trace_guard()** is invoked at the entry site of each basic block for code branch collection.
- the pc_trace_guard_init() assigns ID numbers to each basic block for identification

Coverage Initialization

为什么现有的技术不能被直接用于嵌入式系统？

- for an user land program, **pc_trace_guard_init()** is invoked by a special instrumented function **__sanitizer_cov_module_init()**, which is called before entering each module of the program.
- the variety of compilation environments and the diverse implementation strategies lead to a diverse runtime model and inadequate compilation supports

Coverage Initialization

解决方案:

a **dynamic initialization mechanism, independent of both target architecture and target kernel**

1. the target OS's source code will be compiled into an ELF image.
2. for a function that needs to be invoked, we need to pinpoint the exact position of the function. (**disassemble**)
3. implement an indirect call to it during OS startup through a **function pointer**

OS Agnostic Coverage Collection

现有的coverage collection是如何做的？

Existing coverage collection tools, e.g., **KCOV**, utilizing sophisticated services from target OS, e.g., mmap, fopen, etc.,

解决方案：

collect coverage statistics with several simple binary instruments like bitwise

(不管KCOV这种成熟技术了，自创一些解决方案)

OS Agnostic Coverage Collection

解决思路（基础版）：

To calculate this hit count, we first need to identify which edges a test case covers. Intuitively, we can give each edge a hash number based on its previous and current basic blocks.

存在的问题：

- edge's direction information may loss
- hash操作开销太大

OS Agnostic Coverage Collection

解决思路（升级版）：

- $\text{previous-block-ID} \oplus \text{current-block-ID}$
- shift the previous basic block ID by one bit to the right after execution so that we can distinguish between onward and backward control flow

好处：

- 同时记录了block-ID、direction信息
- xor和shift开销都很小

Coverage Analysis

遇到的问题：

1. 希望在host端完成analysis过程，但是覆盖率数据是在QEMU VM跑的被测试的OS中得到的
2. 希望提高分析效率: coverage comparison is **a frequent operation** that must be performed after each input's execution.

Coverage Analysis

1. **shared coverage buffer**: Tardis proposes using a **shared coverage buffer** to store the coverage information so that it can be accessed on both the guest side and host side directly. (by **shared memory mechenism**)
2. **Efficient Analysis**: a compact bitmap to record the global coverage and check for new recovery.
 - 64KB bitmap (这样就可以放进L2 cache)
 - 对hit count做hash, 每个hit count不超过8bit

(这段没明白, 是对每条边都做还是?)

实验部分

实验环境：

1. UC/OS [13];
2. FreeRTOS [2];
3. Rt-Thread [33]; and
4. Zephyr [15]

The experiments were conducted on a Linux server with 64 GB of memory and a 16-core CPU

实验部分

1. RQ1: Tardis是否能够揭示不同嵌入式操作系统中的新漏洞?
经过测试, 发现了17个previously unknown bugs
2. RQ2: 与黑盒模糊测试相比, Tardis的覆盖率指导机制是否有效地实现了更高的代码覆盖率?
为了测试coverage guide的效率, 用Tardis和裁剪掉coverage guide的Tardis进行对比。
3. RQ3: 在模糊测试期间, 插桩给Tardis带来了哪些开销?
内存: 不同的OS效率不同, 大概10%-40%。时间: 也是30%左右。
测试方法: 选择了不同的测试用例, 测试插桩前和插桩后的运行时间差。

RELATED WORK

kernel fuzzing相关

- syzkaller, 以及自动生成system call description的工作, 例如Healer。
HFL[11] 采用symbolic execution来生成高质量输入来触发新的内核代码执行路径。
- GPOS fuzzing:
 - Gustave,
 - KAFL: 使用Intel-PT来支持coverage-guided fuzzing, 在architecture上有限制
 - Rtkaller: 开源<https://github.com/Rrooach/Rtkaller>。based on Syzkaller, currently only support rt-Linux fuzzing, can not be adapted to other Embedded OS scenarios。

RELATED WORK

Instrumentation in Fuzzing

没有coverage guidance的fuzz效率比较有限，因此收集coverage非常重要。现有的coverage collection方案：

- SanitizerCoverage: SanitizerCoverage offers function level, basic block level, and edge level coverage information
- AFL-GCC: AFL provides edge level coverage only

RELATED WORK

Main Differences between general purpose OS fuzzing and Embedded OS fuzzing

大多数coverage-guided kernel fuzzer需要比较完整的existing infrastructures, such as standard APIs and uniform memory layouts. 但是Embedded OSs可能并不能对fuzzer提供这些支持。

Gustave and KAFL that enable coverage-guided fuzzing of certain Embedded OS by heavily modified QEMU and introducing the binary instrumentation, 但是这样做开销太大。

Thanks!